


(a) NVIDIA Jetson AGX CPU: Quad-core Cortex-A57 + 128 CUDA cores RAM: 4GB LPDDR4 GPU acceleration; significantly faster YOLO inference
(b) Raspberry Pi5 CPU: 2.4GHz Quad-core ARM Cortex-A76 RAM: 16GB LPDDR4 Fastest SBC used; suitable for 5-10FPS inference
(c) Rock Pi 4C+ CPU: 1.8GHz Hexa-core (Cortex-A72 + Cortex-A53) RAM: 4GB LPDDR4 Good thermals; stable inference performance
(d) Raspberry Pi-0 2W CPU: 1.0GHz Quad-core ARM Cortex-A53 RAM: 512MB Extremely constrained; no native Ethernet; slow inference

(e) 5NOK for understanding the edge size used in this work